

# Neural Style Transfer with PyTorch

*In this notebook we will try to merge the worlds of popular series*

## Importing necessary packages

In [1]:

```
!pip3 install torch torchvision
!pip3 install pillow==4.1.1
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.6/dist-packages (1.5.0+cu101)
Requirement already satisfied: torchvision in /usr/local/lib/python3.6/dist-packages (0.6.0+cu101)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from torch) (1.18.5)
Requirement already satisfied: future in /usr/local/lib/python3.6/dist-packages (from torch) (0.16.0)
Requirement already satisfied: pillow>=4.1.1 in /usr/local/lib/python3.6/dist-packages (from torchvision) (4.1.1)
Requirement already satisfied: olefile in /usr/local/lib/python3.6/dist-packages (from pillow>=4.1.1->torchvision) (0.46)
Requirement already satisfied: pillow==4.1.1 in /usr/local/lib/python3.6/dist-packages (4.1.1)
Requirement already satisfied: olefile in /usr/local/lib/python3.6/dist-packages (from pillow==4.1.1) (0.46)
```

In [0]:

```
%matplotlib inline
from PIL import Image

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

import matplotlib.pyplot as plt

import torchvision.transforms as transforms
import torchvision.models as models

import copy
```

## Loading the pictures

In [0]:

```
imsize = 1024

loader = transforms.Compose([
    transforms.Resize(imsize),
    transforms.CenterCrop(imsize),
    transforms.ToTensor()])
```

In [0]:

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
def image_loader(image_name):
    image = Image.open(image_name)
    image = loader(image).unsqueeze(0)
    return image.to(device, torch.float)
```

```
style_img1 = image_loader("/sezam.jpg")
style_img2 = image_loader("/rick.jpg")
content_img = image_loader("/shrek.jpg")
```

## Printing the pictures we will work with

In [62]:

```
unloader = transforms.ToPILImage()



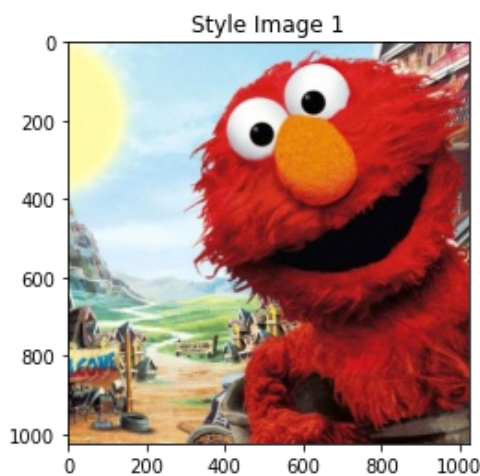
plt.ion()

def imshow(tensor, title=None):
    image = tensor.cpu().clone()
    image = image.squeeze(0)
    image = unloader(image)
    plt.imshow(image)
    if title is not None:
        plt.title(title)
    plt.pause(0.001)

plt.figure()
imshow(style_img1, title='Style Image 1')

plt.figure()
imshow(style_img2, title='Style Image 2')

plt.figure()
imshow(content_img, title='Content Image')
```





## Content Loss

In [0]:

```
class ContentLoss(nn.Module):

    def __init__(self, target,):
        super(ContentLoss, self).__init__()

        self.target = target.detach()
        self.loss = F.mse_loss(self.target, self.target)

    def forward(self, input):
        self.loss = F.mse_loss(input, self.target)
        return input
```

## Gram matrix

In [0]:

```
def gram_matrix(input):
    batch_size, h, w, f_map_num = input.size()

    features = input.view(batch_size * h, w * f_map_num)

    G = torch.mm(features, features.t())

    return G.div(batch_size * h * w * f_map_num)
```

## Style Loss

In [0]:

```
class StyleLoss(nn.Module):
    def __init__(self, target_feature):
        super(StyleLoss, self).__init__()
        self.target = gram_matrix(target_feature).detach()
        self.loss = F.mse_loss(self.target, self.target)

    def forward(self, input):
        G = gram_matrix(input)
        self.loss = F.mse_loss(G, self.target)
        return input
```

## Normalization

In [0]:

```
cnn_normalization_mean = torch.tensor([0.485, 0.456, 0.406]).to(device)
cnn_normalization_std = torch.tensor([0.229, 0.224, 0.225]).to(device)
```

In [0]:

```
class Normalization(nn.Module):
```

```

def __init__(self, mean, std):
    super(Normalization, self).__init__()

    self.mean = torch.tensor(mean).view(-1, 1, 1)
    self.std = torch.tensor(std).view(-1, 1, 1)

def forward(self, img):

    return (img - self.mean) / self.std

```

In [0]:

```

content_layers_default = ['conv_4']
style_layers_default = ['conv_1', 'conv_2', 'conv_3', 'conv_4', 'conv_5']

```

## Pretrained model

In [16]:

```
cnn = models.vgg19(pretrained=True).features.to(device).eval()
```

Downloading: "https://download.pytorch.org/models/vgg19-dcbb9e9d.pth" to /root/.cache/torch/checkpoints/vgg19-dcbb9e9d.pth

## Full model

In [0]:

```

def get_style_model_and_losses(cnn, normalization_mean, normalization_std,
                              content_img, style_img1, style_img2,
                              content_layers=content_layers_default,
                              style_layers=style_layers_default):

    cnn = copy.deepcopy(cnn)

    normalization = Normalization(normalization_mean, normalization_std).to(device)

    content_losses = []
    style_losses1 = []
    style_losses2 = []

    model = nn.Sequential(normalization)

    i = 0
    for layer in cnn.children():
        if isinstance(layer, nn.Conv2d):
            i += 1
            name = 'conv_{}'.format(i)
        elif isinstance(layer, nn.ReLU):
            name = 'relu_{}'.format(i)

            layer = nn.ReLU(inplace=False)
        elif isinstance(layer, nn.MaxPool2d):
            name = 'pool_{}'.format(i)
        elif isinstance(layer, nn.BatchNorm2d):
            name = 'bn_{}'.format(i)
        else:
            raise RuntimeError('Unrecognized layer: {}'.format(layer.__class__.__name__))

    model.add_module(name, layer)

    if name in content_layers:
        target = model(content_img).detach()

```

```

        content_loss = ContentLoss(target)
        model.add_module("content_loss_{}".format(i), content_loss)
        content_losses.append(content_loss)

    if name in style_layers:

        target_feature1 = model(style_img1).detach()
        target_feature2 = model(style_img2).detach()

        style_loss1 = StyleLoss(target_feature1)
        model.add_module("style_loss_1{}".format(i), style_loss1)
        style_losses1.append(style_loss1)

        style_loss2 = StyleLoss(target_feature2)
        model.add_module("style_loss_2{}".format(i), style_loss2)
        style_losses2.append(style_loss2)

    for i in range(len(model) - 1, -1, -1):
        if isinstance(model[i], ContentLoss) or isinstance(model[i], StyleLoss):
            break

    model = model[:i + 1]

    return model, style_losses1, style_losses2, content_losses

```

In [0]:

```

def get_input_optimizer(input_img):

    optimizer = optim.LBFGS([input_img.requires_grad_()])
    return optimizer

```

## Train cycle

In [0]:

```

def run_style_transfer(cnn, normalization_mean, normalization_std,
                      content_img, style_img1, style_img2, input_img, num_steps=500,
                      style_weight=100000, content_weight=1):
    """Run the style transfer."""
    print('Building the style transfer model..')
    model, style_losses1, style_losses2, content_losses = get_style_model_and_losses
(cnn,
    normalization_mean, normalization_std, content_img, style_img1, style_img2)
    optimizer = get_input_optimizer(input_img)

    print('Optimizing..')
    run = [0]
    while run[0] <= num_steps:

        def closure():

            input_img.data.clamp_(0, 1)

            optimizer.zero_grad()

            model(input_img)

            style_score1 = 0
            style_score2 = 0
            content_score = 0

            for sl1 in style_losses1:
                style_score1 += sl1.loss

            for sl2 in style_losses2:
                style_score2 += sl2.loss

            for cl in content_losses:

```

```

        content_score += cl.loss

    style_score1 *= style_weight
    style_score2 *= style_weight
    content_score *= content_weight

    loss = style_score1 + style_score2 + content_score
    loss.backward()

    run[0] += 1
    if run[0] % 50 == 0:
        print("run {}".format(run))
        print('Style Loss 1 : {:.4f} Style Loss 2 : {:.4f} Content Loss: {:.4f}'
              .format(
                  style_score1.item(), style_score2.item(), content_score.item()))
        print()

    return style_score1 + style_score2 + content_score

optimizer.step(closure)

input_img.data.clamp_(0, 1)

return input_img

```

## Training the model on our pictures


In [63]:

```

input_img = content_img.clone()
input_img = torch.randn(content_img.data.size(), device=device)

plt.figure()
imshow(input_img, title='Input Image')
output = run_style_transfer(cnn, cnn_normalization_mean, cnn_normalization_std,
                           content_img, style_img1, style_img2, input_img)

```



Building the style transfer model..  
Optimizing..

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:7: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
import sys
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:8: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

```

```

run [50]:
Style Loss 1 : 142.584015 Style Loss 2 : 96.402061 Content Loss: 17.792370

```

```

run [100]:
Style Loss 1 : 85.896484 Style Loss 2 : 93.578888 Content Loss: 12.928497

```

```
run [150]:  
Style Loss 1 : 84.008652 Style Loss 2 : 89.786438 Content Loss: 10.263264  
  
run [200]:  
Style Loss 1 : 83.289246 Style Loss 2 : 87.740654 Content Loss: 8.654465  
  
run [250]:  
Style Loss 1 : 82.704308 Style Loss 2 : 86.231346 Content Loss: 7.551970  
  
run [300]:  
Style Loss 1 : 82.315414 Style Loss 2 : 84.662270 Content Loss: 6.825212  
  
run [350]:  
Style Loss 1 : 82.333862 Style Loss 2 : 82.370811 Content Loss: 6.552594  
  
run [400]:  
Style Loss 1 : 79.296043 Style Loss 2 : 82.109741 Content Loss: 6.713765  
  
run [450]:  
Style Loss 1 : 78.936890 Style Loss 2 : 79.520493 Content Loss: 6.864116  
  
run [500]:  
Style Loss 1 : 77.228920 Style Loss 2 : 79.087578 Content Loss: 6.282901
```

## Printing the result

In [64]:

```
plt.figure()  
imshow(output, title='Output Image')  
  
plt.ioff()  
plt.show()
```

